Best-First Rao-Blackwellization

J Chiu

March 5, 2022

Gradient Estimation

- Goal: optimize
$$\max_\theta \sum_x p_\theta(x) f(x)$$
(optimizing wrt parameters of f is easy)
- Discrete x, expensive evaluation of $f(x)$

Gradient Estimators

- Score-function estimator (SFE)

$$\nabla_\theta \sum_x p_\theta(x) f(x) = \mathbb{E}_{p_\theta(x)} \left[f(x) \nabla \log p_\theta(x) \right]$$

  - High variance, consistent
- Continuous relaxation, $x \approx g(\theta, \epsilon)$

$$\nabla_\theta \sum_x p_\theta(x) f(x) \approx \mathbb{E}_{p(\epsilon)} \left[\nabla_\theta f(g_\theta(\theta, \epsilon)) \right]$$

  - Biased, lower variance
  - Requires relaxable f
  - Stochastic softmax tricks (SST)
- Focus on improving SFE
  - Total error most important, but SFE under-explored

Rao-Blackwellization of score function estimator

- Condition on values of $f(x)$ for $x \in S$

$$
\begin{aligned}
&\nabla_\theta \sum_x p_\theta(x) f(x) \\
&= \mathbb{E}_{p_\theta(x)}\left[f(x) \nabla \log p_\theta(x)\right] \\
&= \sum_{x \in S} p(x) f(x) \nabla \log p_\theta(x) + \mathbb{E}_{p_\theta(x \notin S)}\left[f(x) \nabla \log p_\theta(x)\right]
\end{aligned}
$$

- Conditioning on all values $f(x)$ = no variance
- Want S to have x with high $p(x)f(x)$
  - Need to compromise with high $p(x)$
  - Use heuristic estimate of $f(x)$?
  - What if x is structured?

Structured Setting

Flat

p(1)	p(2)	p(3)	p(4)
f(1)	f(2)	f(3)	f(4)

Structured

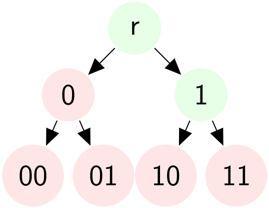

Best-first Rao-Blackwellization

- Modern MCTS
  - Prior $p(x)$ (forward)
  - Cost-to-go estimate $\tilde{V}(x)$ (backward)
  - Search procedure that links the two
- Proposal
  - In many cases, already have the prior $p(x)$
  - Estimate cost-to-go with continuous relaxation or cheaper problem-dependent weaker model (Markov transformer)
  -

BF RB

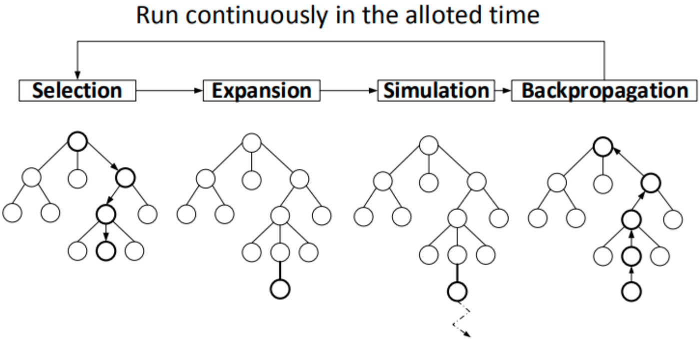
Run continuously in the alloted time
Selection
Expansion
Simulation
Backpropagation

Citations I